



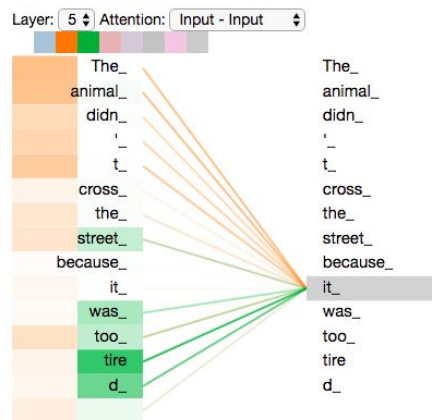
FNet

Mixing Tokens with Fourier Transform



What it's all about

- Transformer magic lies in Self-Attention
- Self-attention is expensive
 - Quadratic cost, $O(N^2)$ where N is the size of the sequence
 - (Remember perceiver)



What's the attention of "it"?

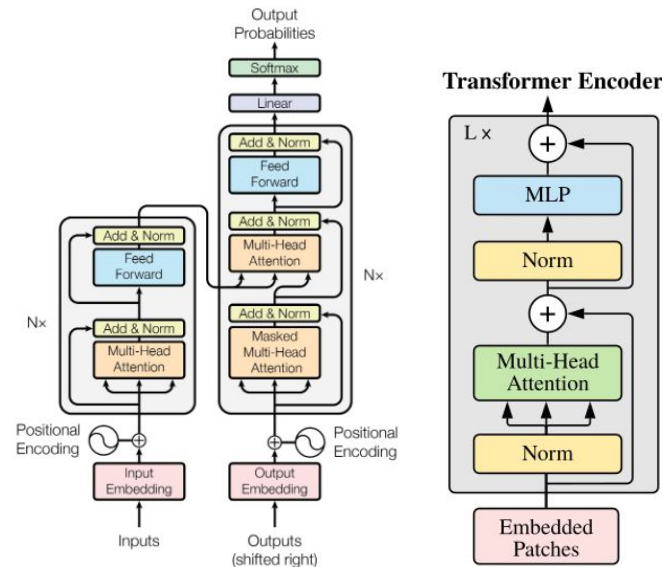


Figure 1: The Transformer - model architecture.

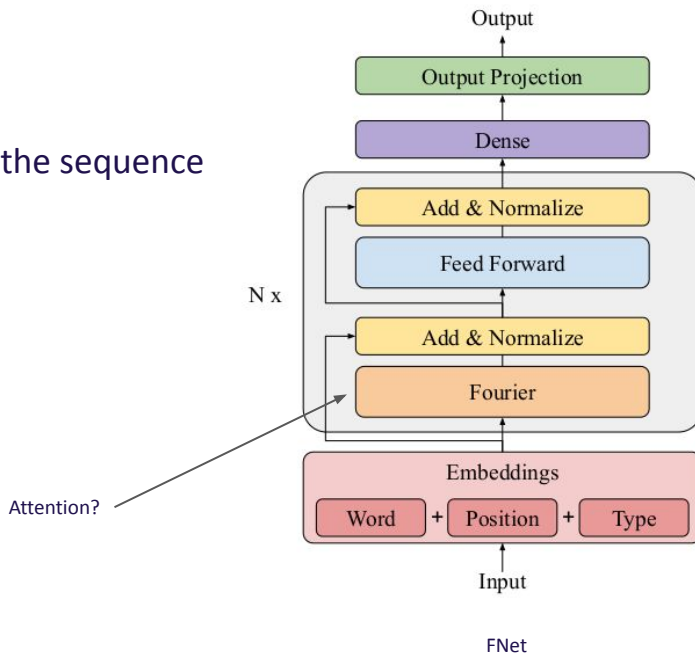
(from Vaswani et. al Attention is all you need)

(from Dosovitskiy et. al An image is worth 16x16 words.)



What it's all about

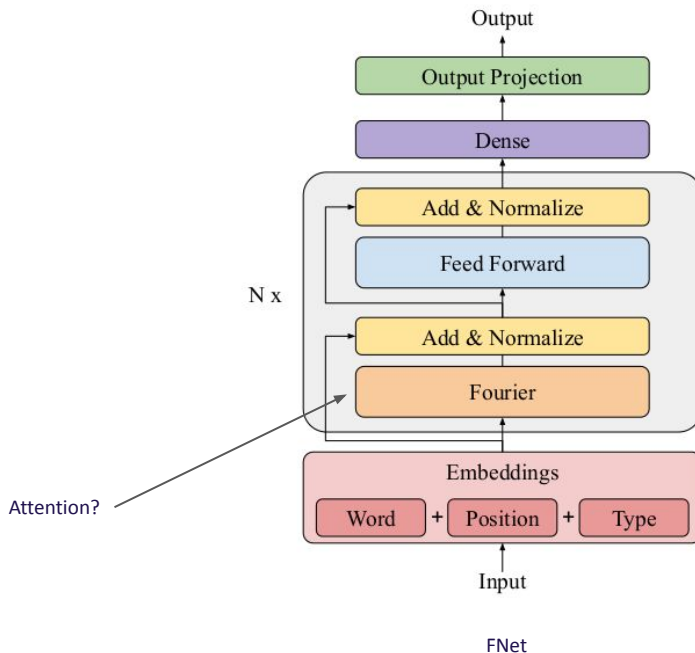
- Transformer magic lies in Self-Attention
- Self-attention is expensive
 - Quadratic cost, $O(N^2)$ where N is the size of the sequence
 - (Remember perceiver)
- Can we do better?
 - Kind of





Fourier Transform? (I)

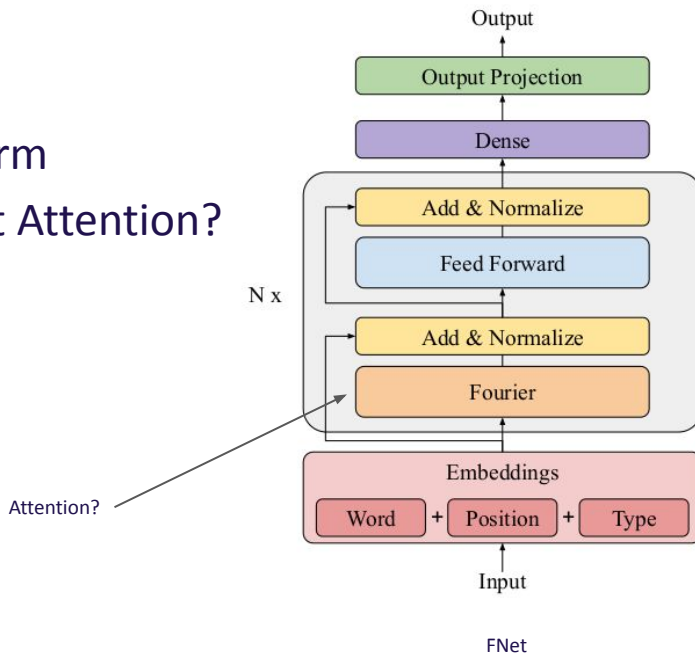
- Where is the learning part?
 - No parameters
- Does it work?
 - Yes
- Is it more efficient?
 - A lot! (obviously, no attention cost)
- Is FT all we need?
 - No





Fourier Transform? (II)

- Fourier transform works better
 - but that's not the point, probably
- Fast implementations of Fourier Transform
- The real question is: Can we live without Attention?
 - We'll see





What's in the box?! (Fourier Version)

- How does this translates to our problem?
 - We don't have signals as input
- In this case we have a sequence of tokens
 - and tokens have embeddings
- 2 Fourier transforms.
 - One sequence wise
 - One along the hidden dimensions

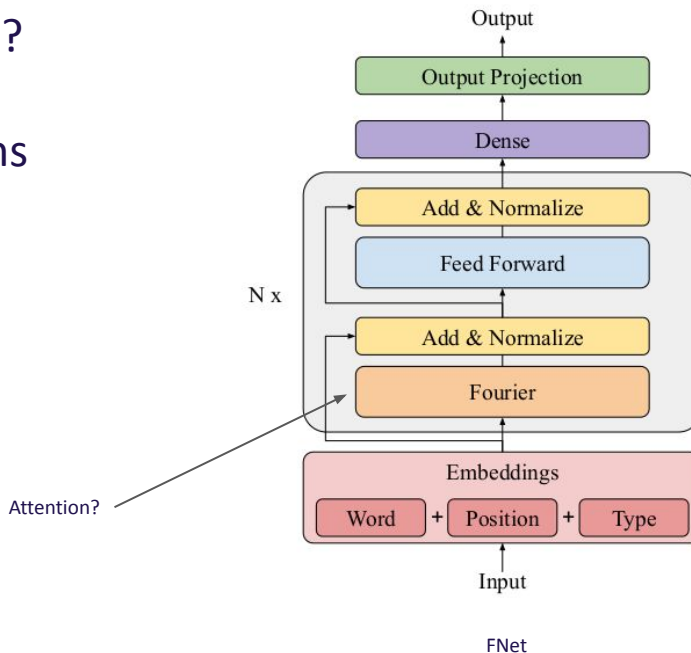
$$y = \Re(\mathcal{F}_{\text{seq}}(\mathcal{F}_h(x))).$$

Input

$$X_k = \sum_{n=0}^{N-1} x_n \cdot e^{-\frac{i2\pi}{N} kn}$$

$$= \sum_{n=0}^{N-1} x_n \cdot \left[\cos\left(\frac{2\pi}{N} kn\right) - i \cdot \sin\left(\frac{2\pi}{N} kn\right) \right], \quad (\text{Eq.1})$$

Discrete fourier transform reminder. (We drop the imaginary part)

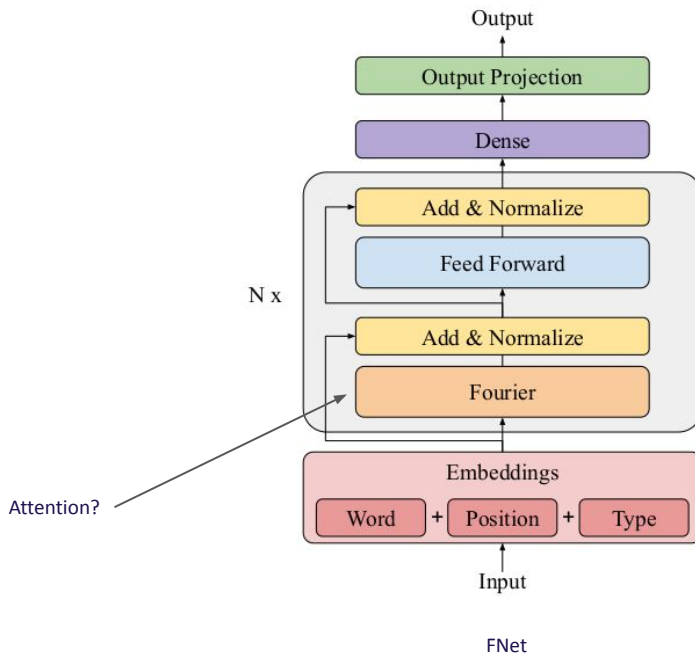





What's in the box?! (Fourier Version)

- 2 Fourier transforms.
 - One sequence wise
 - One along the hidden dimensions

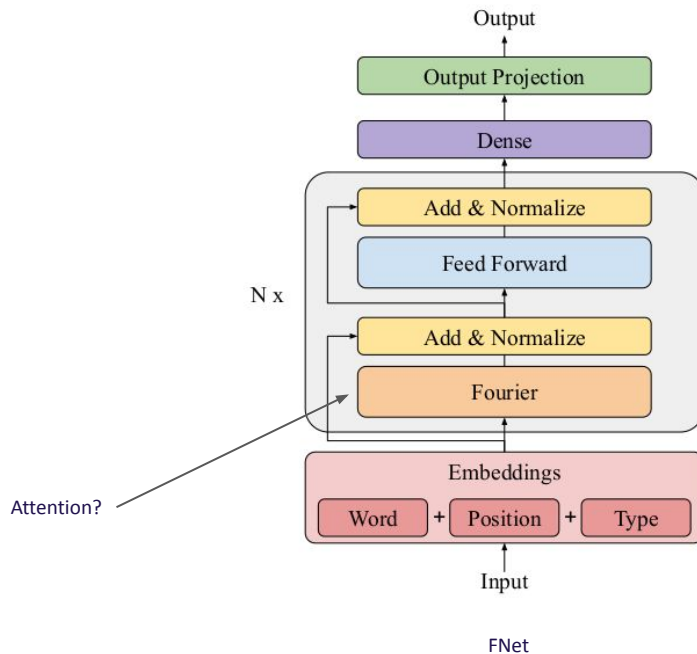
$$y = \Re(\mathcal{F}_{\text{seq}}(\mathcal{F}_h(x))).$$





Token mixing

- The success of the method is most likely caused by the mixing of tokens.
- Attention is a fancier way of doing that by learning what matters for what.



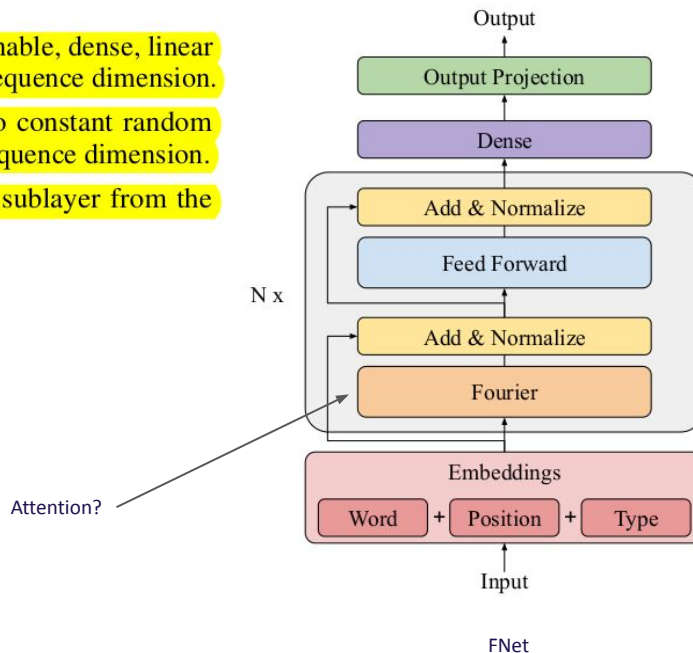


Token mixing - Other flavours

- Linear encoder: we replace each self-attention sublayer with a two learnable, dense, linear sublayers, one applied to the hidden dimension and one applied to the sequence dimension.
- Random encoder: we replace each self-attention sublayer with a two constant random matrices, one applied to the hidden dimension and one applied to the sequence dimension.
- Feed Forward-only (FF-only) encoder: we remove the self-attention sublayer from the Transformer layers; this model has no token mixing.

Hey bro!

MLP-Mixer: An all-MLP Architecture for Vision





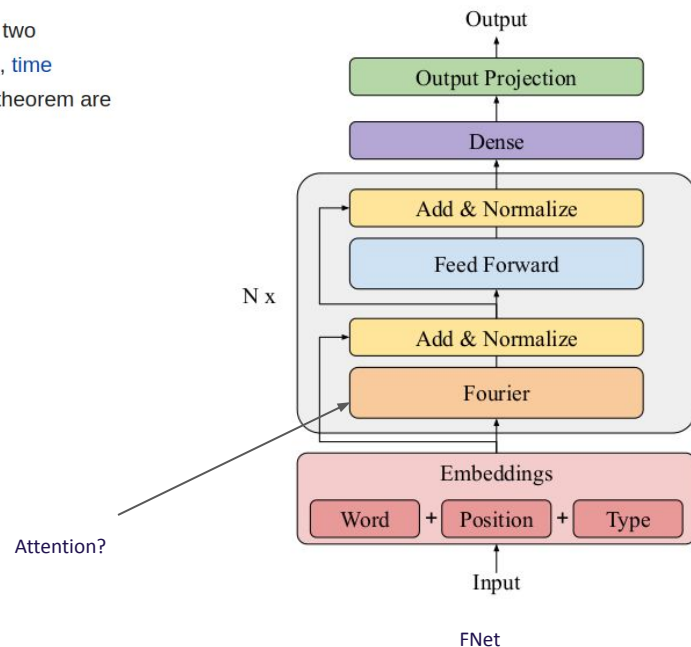
Another all MLP network?

Convolution theorem

From Wikipedia, the free encyclopedia

In [mathematics](#), the **convolution theorem** states that under suitable conditions the [Fourier transform](#) of a [convolution](#) of two functions (or [signals](#)) is the [pointwise product](#) of their Fourier transforms. More generally, convolution in one domain (e.g., [time domain](#)) equals point-wise multiplication in the other domain (e.g., [frequency domain](#)). Other versions of the convolution theorem are applicable to various [Fourier-related transforms](#).

- In our case we constantly move from one domain to another in every layer →
- Convolutions kind of exist in FNet





Experiments

Table 1: GLUE Validation results on TPUs, after finetuning on respective tasks. We report the mean of accuracy and F1 scores for QQP and MRPC, Spearman correlations for STS-B and accuracy scores for all other tasks. The MNLI metrics are reported by the match/mismatch splits. Average scores exclude any failure cases. After controlling for batch size and training steps, the GPU metrics (not shown) are similar.

Model	MNLI	QQP	QNLI	SST-2	CoLA	STS-B	MRPC	RTE	Avg.
BERT-Base	84/81	87	91	93	73	89	83	69	83.3
Linear-Base	74/75	84	80	94	67	67	83	69	77.0
FNet-Base	72/73	83	80	95	69	79	76	63	76.7
Random-Base	51/50	70	61	76	67	4	73	57	56.6
FF-only-Base	34/35	31	52	48	67	FAIL	73	54	49.3
FNet-Hybrid-Base	78/79	85	88	94	76	86	79	60	80.6
BERT-Large	88/88	88	92	95	71	88	86	66	84.7
Linear-Large	35/36	84	80	79	67	24	73	60	59.8
FNet-Large	78/76	85	85	94	78	84	88	69	81.9

Bye

